

# LOOP OPTIMIZATIONS WITH OPENACC

Speaker



# LECTURE 3 OUTLINE


## Topics to be covered

- Gangs, Workers, and Vectors Demystified
- GPU Profiles
- Loop Optimizations
- Week 3 Lab
- Where to Get Help

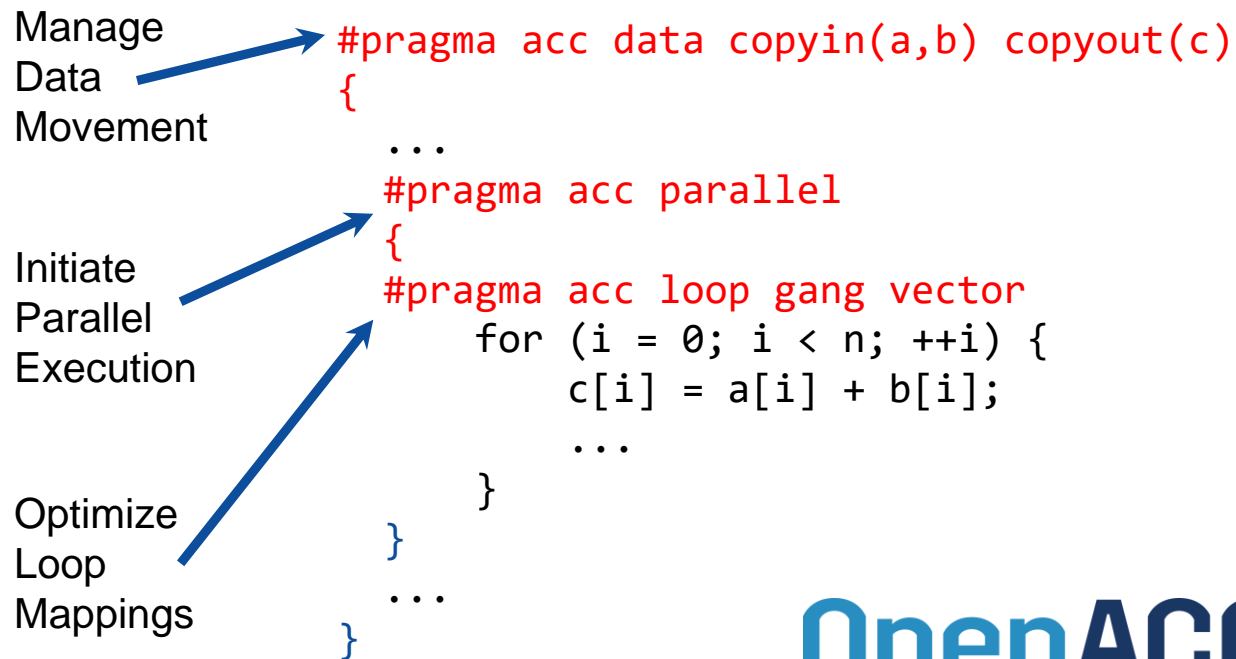
# WEEKS 1 & 2 REVIEW

# OPENACC DEVELOPMENT CYCLE

- **Analyze** your code to determine most likely places needing parallelization or optimization.
- **Parallelize** your code by starting with the most time consuming parts and check for correctness.
- **Optimize** your code to improve observed speed-up from parallelization.



# OpenACC Directives



- Incremental
- Single source
- Interoperable
- Performance portable
- CPU, GPU, Manycore

**OpenACC**  
Directives for Accelerators

# PARALLELIZE WITH OPENACC PARALLEL LOOP

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
  
    #pragma acc parallel loop reduction(max:err)  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
  
    #pragma acc parallel loop  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
    iter++;  
}
```

Parallelize first loop nest,  
max *reduction* required.

Parallelize second loop.

We didn't detail *how* to  
parallelize the loops, just *which*  
loops to parallelize.

# OPTIMIZED DATA MOVEMENT

```
#pragma acc data copy(A[:n*m]) copyin(Anew[:n*m])
while ( err > tol && iter < iter_max ) {
    err=0.0;

    #pragma acc parallel loop reduction(max:err) copyin(A[0:n*m])
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {

            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);


            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }

    #pragma acc parallel loop copyin(Anew[0:n*m]) copyout(A[0:n*m])
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }
    iter++;
}
```

Copy A to/from the accelerator only when needed.

Copy initial condition of Anew, but not final value

# OPENACC SPEED-UP





# GANGS, WORKERS, AND VECTORS DEMYSTIFIED

# GANGS, WORKERS, AND VECTORS DEMYSTIFIED



# GANGS, WORKERS, AND VECTORS DEMYSTIFIED




# GANGS, WORKERS, AND VECTORS DEMYSTIFIED

!




How much work 1 worker  
can do is limited by his  
speed.

A single worker can only  
move so fast.





# GANGS, WORKERS, AND VECTORS DEMYSTIFIED



Even if we increase the size of his roller, he can only paint so fast.

We need more workers!



# GANGS, WORKERS, AND VECTORS DEMYSTIFIED



Multiple workers can do more work and share resources, if organized properly.



# GANGS, WORKERS, AND VECTORS DEMYSTIFIED

By organizing our workers into groups (gangs), they can effectively work together within a floor.

Groups (gangs) on different floors can operate independently.

Since gangs operate independently, we can use as many or few as we need.



# GANGS, WORKERS, AND VECTORS DEMYSTIFIED

Even if there's not enough gangs for each floor, they can move to another floor when ready.



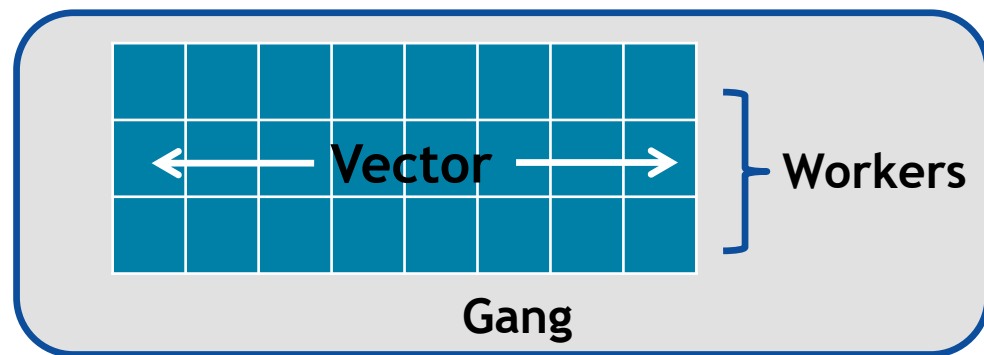


# GANGS, WORKERS, AND VECTORS

Our painter is like an OpenACC **worker**, he can only do so much.

His roller is like a **vector**, he can move faster by covering more wall at once.

Eventually we need more workers, which can be organized into **gangs** to get more done.

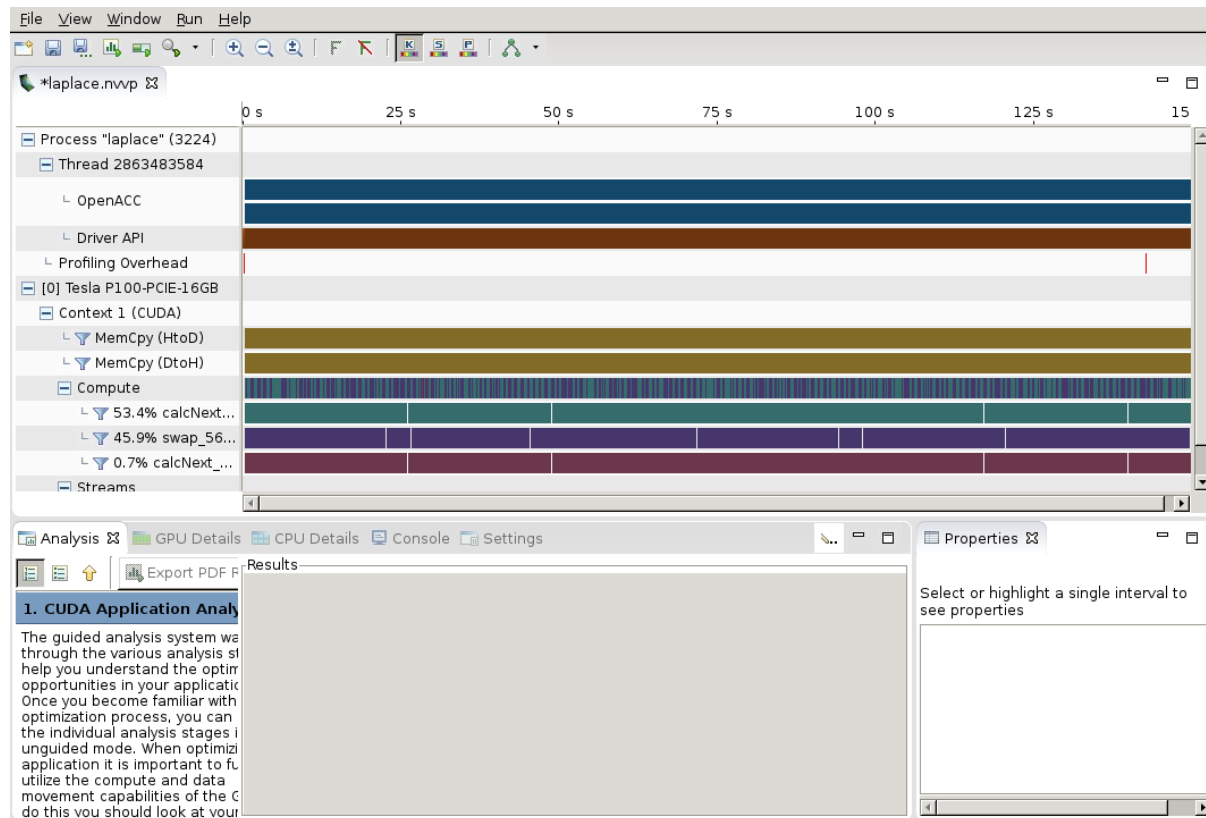


# GPU PROFILES

# PROFILING GPU CODE (PGPROF)

## Using PGPROF to profile GPU code


- PGPROF presents far more information when running on a GPU
- We can view CPU Details, GPU Details, a Timeline, and even do Analysis of the performance



# PROFILING GPU CODE (PGPROF)

## Using PGPROF to profile GPU code

- **MemCpy(HtoD):** This includes data transfers from the Host to the Device (CPU to GPU)
- **MemCpy(DtoH):** These are data transfers from the Device to the Host (GPU to CPU)
- **Compute:** These are our computational functions. We can see our calcNext and swap function

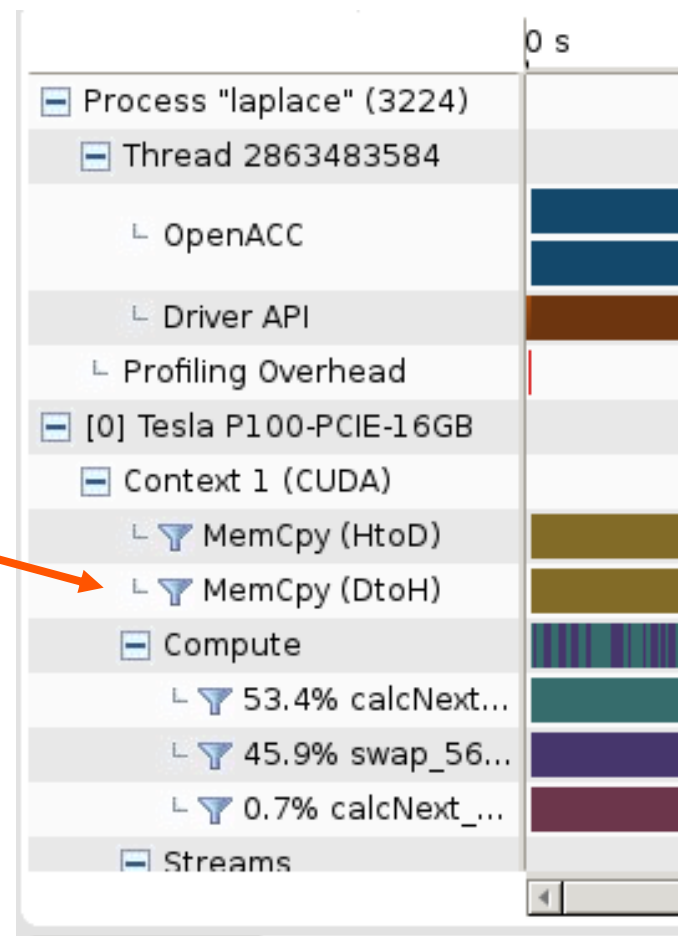




# PROFILING GPU CODE (PGPROF)

## Using PGPROF to profile GPU code


- **MemCpy(HtoD):** This includes data transfers from the Host to the Device (CPU to GPU)
- **MemCpy(DtoH):** These are data transfers from the Device to the Host (GPU to CPU)
- **Compute:** These are our computational functions. We can see our calcNext and swap function



# PROFILING GPU CODE (PGPROF)

## Using PGPROF to profile GPU code

- **MemCpy(HtoD):** This includes data transfers from the Host to the Device (CPU to GPU)
- **MemCpy(DtoH):** These are data transfers from the Device to the Host (GPU to CPU)
- **Compute:** These are our computational functions. We can see our calcNext and swap function



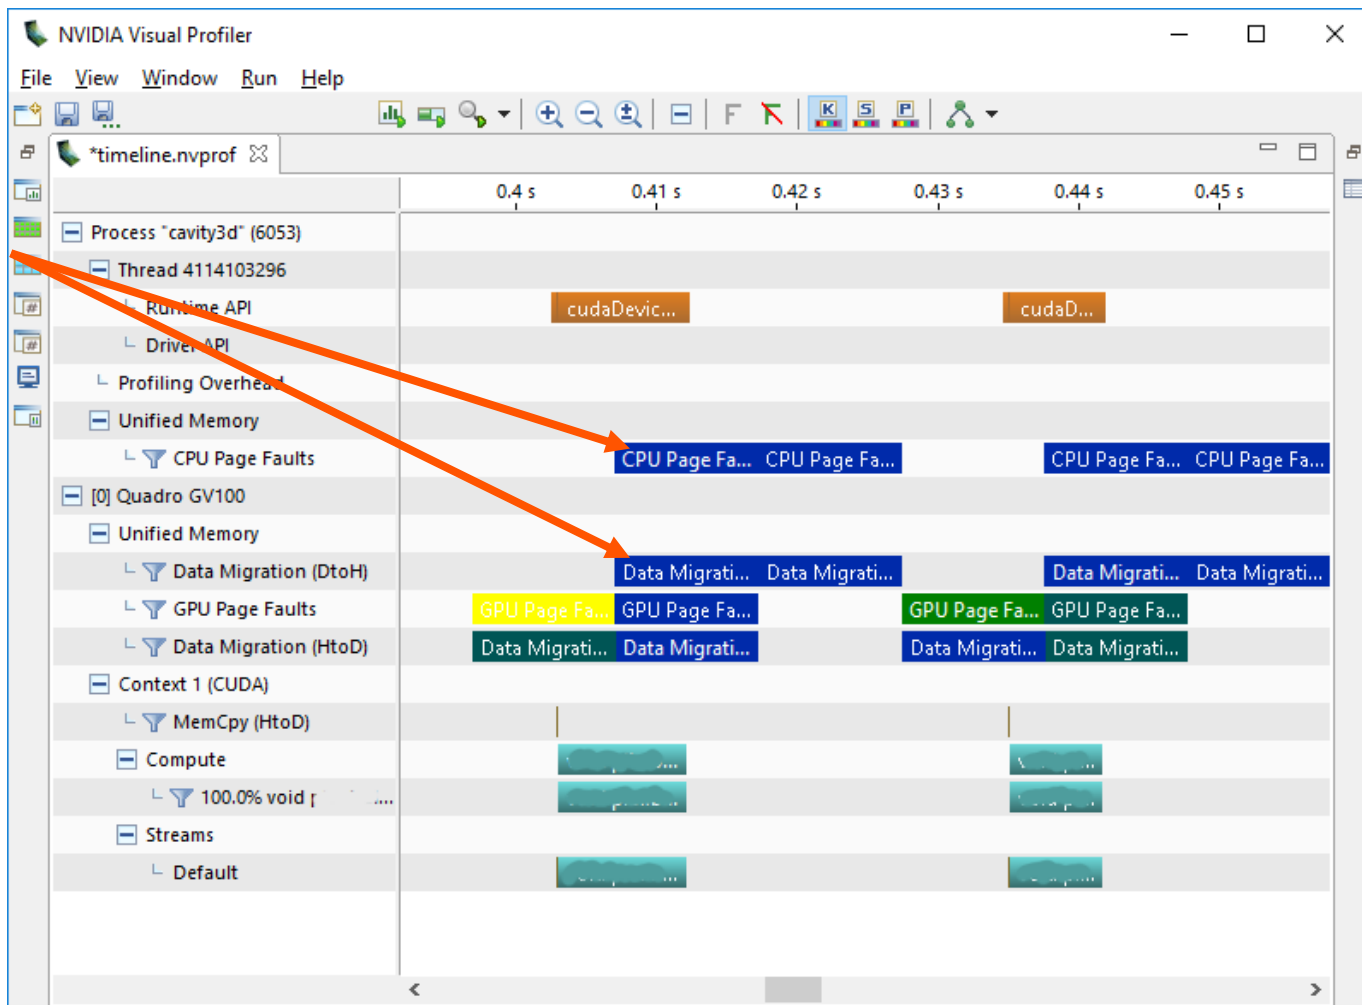
# PROFILING GPU CODE

## Managed Memory

CPU Page Faults trigger Device to Host Migrations.

GPU Page Faults trigger Host to Device Migrations.

These may be hints that more needs to be parallelized and/or data optimization is needed.




# PROFILING GPU CODE

## Managed Memory

CPU Page Faults trigger Device to Host Migrations.

GPU Page Faults trigger Host to Device Migrations.

These may be hints that more needs to be parallelized and/or data optimization is needed.





# LOOP OPTIMIZATIONS

# OPENACC LOOP DIRECTIVE

## Expressing parallelism

- Mark a single for loop for parallelization
- Allows the programmer to give additional information and/or optimizations about the loop
- Provides many different ways to describe the type of parallelism to apply to the loop
- Must be contained within an OpenACC compute region (either a kernels or a parallel region) to parallelize loops

### C/C++

```
#pragma acc loop  
for(int i = 0; i < N; i++)  
    // Do something
```

### Fortran

```
!$acc loop  
do i = 1, N  
    ! Do something
```

# COLLAPSE CLAUSE

- **collapse( N )**
- Combine the next N tightly nested loops
- Can turn a multidimensional loop nest into a single-dimension loop
- This can be extremely useful for increasing memory locality, as well as creating larger loops to expose more parallelism

```
#pragma acc parallel loop collapse(2)
for( i = 0; i < size; i++ )
    for( j = 0; j < size; j++ )
        double tmp = 0.0f;
        #pragma acc loop reduction(+:tmp)
        for( k = 0; k < size; k++ )
            tmp += a[i][k] * b[k][j];
        c[i][j] = tmp;
```

# COLLAPSE CLAUSE

**collapse( 2 )**

(0,0)	(0,1)	(0,2)	(0,3)
(1,0)	(1,1)	(1,2)	(1,3)
(2,0)	(2,1)	(2,2)	(2,3)
(3,0)	(3,1)	(3,2)	(3,3)

```
#pragma acc parallel loop collapse(2)
for( i = 0; i < 4; i++ )
    for( j = 0; j < 4; j++ )
        array[i][j] = 0.0f;
```

# COLLAPSE CLAUSE

## When/Why to use it

- A single loop might not have enough iterations to parallelize
- Collapsing outer loops gives more scalable (gangs) parallelism
- Collapsing inner loops gives more tight (vector) parallelism
- Collapsing all loops gives the compiler total freedom, but may cost data locality

# COLLAPSE CLAUSE

```
#pragma acc data copy(A[:n*m]) copyin(Anew[:n*m])
while ( err > tol && iter < iter_max ) {
    err=0.0;
```


```
#pragma acc parallel loop reduction(max:err) collapse(2) \
    copyin(A[0:n*m]) copy(Anew[0:n*m])
for( int j = 1; j < n-1; j++) {
    for(int i = 1; i < m-1; i++) {
        Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                           A[j-1][i] + A[j+1][i]);
        err = max(err, abs(Anew[j][i] - A[j][i]));
    }
}
```

```
#pragma acc parallel loop collapse(2) \
    copyin(Anew[0:n*m]) copyout(A[0:n*m])
for( int j = 1; j < n-1; j++) {
    for( int i = 1; i < m-1; i++ ) {
        A[j][i] = Anew[j][i];
    }
}
iter++;
```

Collapse 2 loops into one  
for more flexibility in  
parallelizing.



# OPENACC SPEED-UP



# TILE CLAUSE


- **tile ( x , y , z , ... )**
- Breaks multidimensional loops into “tiles” or “blocks”
- Can increase data locality in some codes
- Will be able to execute multiple “tiles” simultaneously

```
#pragma acc kernels loop tile(32, 32)
for( i = 0; i < size; i++ )
    for( j = 0; j < size; j++ )
        for( k = 0; k < size; k++ )
            c[i][j] += a[i][k] * b[k][j];
```

# TILE CLAUSE

```
#pragma acc kernels loop tile(2,2)
for(int x = 0; x < 4; x++){
    for(int y = 0; y < 4; y++){
        array[x][y]++;
    }
}
```

tile ( 2 , 2 )



# OPTIMIZED DATA MOVEMENT

```
#pragma acc data copy(A[:n*m]) copyin(Anew[:n*m])
while ( err > tol && iter < iter_max ) {
    err=0.0;

    #pragma acc parallel loop reduction(max:err) tile(32,32) \
        copyin(A[0:n*m]) copy(Anew[0:n*m])
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);
            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }

    #pragma acc parallel loop tile(32,32) \
        copyin(Anew[0:n*m]) copyout(A[0:n*m])
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }
    iter++;
}
```

Create 32x32 tiles of the loops to better exploit data locality.

# TILING RESULTS (V100)


The collapse clause often requires an exhaustive search of options.

For our example code...

- CPU saw no benefit from tiling
- GPU saw anywhere from a 23% loss of performance to a 10% improvement

	CPU Improvement	GPU Improvement
Baseline	1.00X	1.00X
4x4	1.00X	0.77X
4x8	1.00X	0.92X
8x4	1.00X	0.95X
8x8	1.00X	1.02X
8x16	1.00X	0.99X
16x8	1.00X	1.02X
16x16	1.00X	1.01X
16x32	1.00X	1.06X
32x16	1.00X	1.07X
32x32	1.00X	1.10X

# OPENACC SPEED-UP





# GANG, WORKER, AND VECTOR CLAUSES

- The developer can instruct the compiler which levels of parallelism to use on given loops by adding clauses:
- **gang** – Mark this loop for gang parallelism
- **worker** – Mark this loop for worker parallelism
- **vector** – Mark this loop for vector parallelism

These can be combined on the same loop.

```
#pragma acc parallel loop gang
for( i = 0; i < size; i++ )
    #pragma acc loop worker
    for( j = 0; j < size; j++ )
        #pragma acc loop vector
        for( k = 0; k < size; k++ )
            c[i][j] += a[i][k] * b[k][j];
```

```
#pragma acc parallel loop \
    collapse(3) gang vector
for( i = 0; i < size; i++ )
    for( j = 0; j < size; j++ )
        for( k = 0; k < size; k++ )
            c[i][j] += a[i][k] * b[k][j];
```

# SEQ CLAUSE

- The **seq** clause (short for sequential) will tell the compiler to run the loop sequentially
- In the sample code, the compiler will parallelize the outer loops across the parallel threads, but each thread will run the inner-most loop sequentially
- The compiler may automatically apply the seq clause to loops as well

```
#pragma acc parallel loop
for( i = 0; i < size; i++ )
    #pragma acc loop
    for( j = 0; j < size; j++ )
        #pragma acc loop seq
        for( k = 0; k < size; k++ )
            c[i][j] += a[i][k] * b[k][j];
```

# PRIVATE AND FIRSTPRIVATE CLAUSES

- The **private** clause allows the programmer to define a list of variables as “thread-private”.
- Each thread will be given a private copy of every variable in the comma-separated list
- **firstprivate** is like private except that the private values are initialized to the same value used on the host. **private** variables are uninitialized.

```
double tmp[3];

#pragma acc kernels loop private(tmp[0:3])
for( i = 0; i < size; i++ )
{
    tmp[0] = <value>;
    tmp[1] = <value>;
    tmp[2] = <value>;
}

// note that the host value of “tmp”
// remains unchanged.
```

# PRIVATE AND FIRSTPRIVATE CLAUSES

- Variables in **private** or **firstprivate** clause are private to the loop level on which the clause appears.
- Private variables on an outer loop are shared within inner loops.

```
double tmp[3];

#pragma acc kernels loop private(tmp[0:3])
for( i = 0; i < size; i++ ) {
    // the tmp array is private to each iteration
    // of the outer loop
    tmp[0] = <value>;
    tmp[1] = <value>;
    tmp[2] = <value>;
    #pragma acc loop
    for ( j = 0; j < size2; j++ ) {
        // but tmp is shared amongst the threads
        // in the inner loop
        array[i][j] = tmp[0]+tmp[1]+tmp[2];
    }
}
```

# SCALARS AND PRIVATE CLAUSE

- By default, scalars are **firstprivate** when used in a parallel region and **private** when used in a kernels region.
- Except in some cases, scalars do not need to be added to a private clause. These cases may include but are not limited to:
  1. Scalars with global storage such as global variables in C/C++, Module variables in Fortran
  2. When the scalar is passed by reference to a device subroutine
  3. When the scalar is used as an rvalue after the compute region, aka “live-out”
- Note that putting scalars in a private clause may actually hurt performance!

# ADJUSTING GANGS, WORKERS, AND VECTORS

The compiler will choose a number of gangs, workers, and a vector length for you, but you can change it with clauses.

- **num\_gangs(N)** – Generate N gangs for this parallel region
- **num\_workers(M)** – Generate M workers for this parallel region
- **vector\_length(Q)** – Use a vector length of Q for this parallel region

```
#pragma acc parallel num_gangs(2) \
    num_workers(2) vector_length(32)
{
    #pragma acc loop gang worker
    for(int x = 0; x < 4; x++){
        #pragma acc loop vector
        for(int y = 0; y < 32; y++){
            array[x][y]++;
        }
    }
}
```




# COLLAPSE CLAUSE WITH VECTOR LENGTH

```
#pragma acc data copy(A[:n*m]) copyin(Anew[:n*m])
while ( err > tol && iter < iter_max ) {
    err=0.0;

    #pragma acc parallel loop reduction(max:err) collapse(2) vector_length(1024) \
        copyin(A[0:n*m]) copy(Anew[0:n*m])
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);
            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }

    #pragma acc parallel loop collapse(2) vector_length(1024) \
        copyin(Anew[0:n*m]) copyout(A[0:n*m])
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }
    iter++;
}
```

# OPENACC SPEED-UP



# LOOP OPTIMIZATION RULES OF THUMB

- It is rarely a good idea to set the number of gangs in your code, let the compiler decide.
- Most of the time you can effectively tune a loop nest by adjusting only the vector length.
- It is rare to use a worker loop. When the vector length is very short, a worker loop can increase the parallelism in your gang.
- When possible, the vector loop should step through your arrays
- Gangs should come from outer loops, vectors from inner

# CLOSING REMARKS

# KEY CONCEPTS

In this lab we discussed...

- Some details that are available to use from a GPU profile
- Gangs, Workers, and Vectors Demystified
- Collapse clause
- Tile clause
- Gang/Worker/Vector clauses

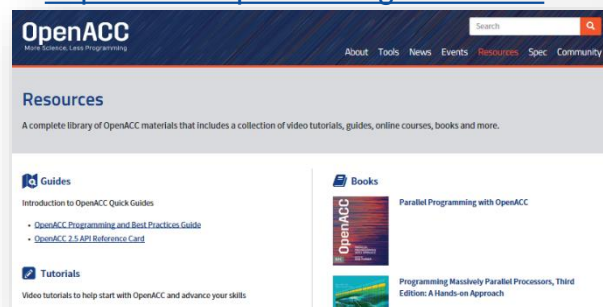
# OPENACC RESOURCES

Guides • Talks • Tutorials • Videos • Books • Spec • Code Samples • Teaching Materials • Events • Success Stories • Courses • Slack • Stack Overflow



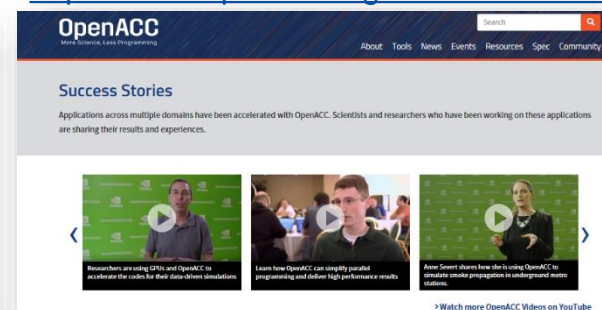
## Resources

<https://www.openacc.org/resources>



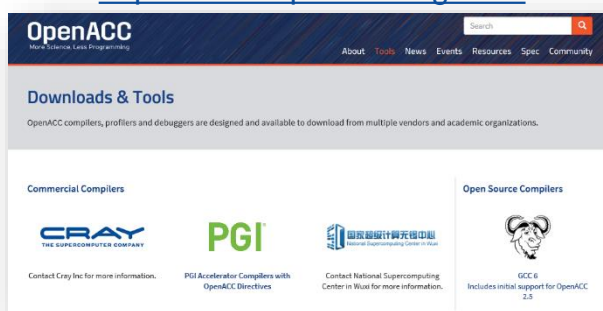
## Success Stories

<https://www.openacc.org/success-stories>



## Compilers and Tools

<https://www.openacc.org/tools>



## Events

<https://www.openacc.org/events>

